Assessment follow-ups

Questions asked to the client during the assignment, the answers, and what each one changes in the build plan (docs/plan.md). These answers settle points the brief leaves open; where they differ from the brief, they win.

Access and privacy (step 3)

ASKED	How is a user linked to "their" products? (proposed: product categories in a config file)
ANSWER	By **brand**, not category. Each manager manages one or more brands and analyzes only those; the CEO sees everything. The mapping isn't in the dataset, so a config file or access table is fine. Customer and order data is visible only through purchases of in-scope products.
PLAN	A users file maps each user to brands (CEO: all), picked at start with --user. Scoping is enforced in the gateway, not the prompt: orders and users are reachable only through order items of in-scope products.

ASKED	Which fields count as PII?
ANSWER	Never shown: first and last name, email, street address, postal code, exact coordinates, IP address. City, state, country, age, gender and traffic source only in aggregates. PII must **never reach the LLM**, so mask it before data goes to the model; keyword or column-name matching alone is not robust enough.
PLAN	PII and aggregate-only columns are tagged in schema.yaml and resolved from the parsed query, so aliases and expressions can't hide them. Results are checked again before rows reach the model. IP address is listed as PII in the HLD.

ASKED	May individual customers appear in results, e.g. "top 10 customers"?
ANSWER	Yes, by user_id or a stable pseudonymous ID (e.g. a salted hash), so follow-ups still work.
PLAN	user_id is replaced by a salted-hash pseudonym that stays the same across turns.

ASKED	Should small groups in aggregates be hidden? (a group can still be one person; proposed minimum: 5)
ANSWER	Yes, hide them. 5 is a good default; keep it configurable and note it in the design doc.
PLAN	MIN_GROUP_SIZE setting, default 5. How the gateway counts customers behind a row is a step 3 design choice.

ASKED	Does IP address matter for the prototype? (it appears only in the events table)
ANSWER	If the prototype doesn't touch events, there's nothing to handle. List it as PII in the design.
PLAN	No code change: the guard already rejects queries on events. Done in the HLD.

Saved reports (step 5)

ASKED	Soft or hard delete? (proposed: soft delete with a confirmation listing what goes)
ANSWER	**Hard delete**, no recovery window. Keep the strict confirmation that shows the exact list of reports first. Only the report's author can see or delete it.
PLAN	Code, not the model, shows the exact list and asks for typed confirmation (a LangGraph interrupt), then deletes for good. Ownership is checked in the store query.

ASKED	What does a saved report contain? (proposed: the question, the SQL and the answer, visible only to its author)
ANSWER	Question, SQL and answer is right; private to the author.
PLAN	The report store keeps those three, plus owner, conversation id and creation time, so both deletion examples from the brief work.

LLM and cost (steps 1, 4 and 7)

ASKED	Which provider will the client's team use to run the prototype?
ANSWER	Gemini. The team tests with its own **free Google AI Studio key**, so stay within free-tier rate limits at roughly **2-5 LLM calls per question**. A sensible default is to fall back between Gemini models only, with other providers in config but switched off. A swappable provider is a nice touch; explain the trade-off in the design.
PLAN	Built: Gemini is the default, a question is capped at 5 LLM calls (a normal one uses 2), and a failed call falls back to another Gemini model (LLM_FALLBACK_MODELS). Other providers stay switchable in config but off by default (Claude is an optional install extra). The trade-off is in docs/model-switching.html and goes into the technical write-up (step 7).

Definitions (step 1)

ASKED	How is revenue defined? (proposed: sum of sale_price, excluding cancelled and returned items)
ANSWER	A sensible default. State the definition in the answer and let users override it.
PLAN	Built: the prompt makes the agent state the definition it used and accept the user's own.

Open

A slot for the presentation call next week has been requested (Monday or Tuesday evening CET preferred); no reply yet. It doesn't change the plan.